

CS 1653: Applied Cryptography and Network Security

Fall 2023

Term Project, Phase 3

Assigned: Wed, Nov 01

Due: Sun, Nov 19 11:59 PM

1 Background

At this point in the semester, your group has developed a fully-functional distributed system; congratulations! This system should currently be capable of storing and retrieving resources from a collection of distributed servers, but should *not* yet offer any protection from clients or servers that behave maliciously. In this phase of the project, we will begin to harden your system against a variety of common security threats. Later sections of this assignment detail a threat model that describes the types of assumptions that you are permitted to make regarding the behavior and intentions of each class of principals in the system. A list of *specific* classes of threats for which your system must provide protections is then described in detail.

Your deliverables for this phase of the project will include (i) a writeup describing your proposed protections for each type of threat, as well as a description of why these protections are sufficient, and (ii) a set of modified applications that implement your protections. Your group's grade for this project will be based on the correctness of the protections described in your preliminary writeup, the accuracy with which your implementation embodies these protections, and a live demonstration of your system.

2 Trust Model

In this phase of the project, we are going to focus on implementing a subset of the security features that will be required of our trustworthy distributed system. Prior to describing the specific threats for which you must provide protections, we now characterize the behavior of the four classes of principals that may be present in our system:

- **Authentication Server** The authentication server is entirely trustworthy. In this phase of the project, this means that the authentication server will only issue tokens to *properly authenticated* clients and will properly enforce the constraints on group creation, deletion, and management specified in the previous phase of the project.
- **Resource Servers** In this phase of the project, you may assume that *properly authenticated* resource servers are entirely trustworthy. In particular, you do not need to worry about a properly authenticated resource server corrupting files, leaking files to unauthorized users, or stealing user tokens. (However, another entity may attempt to impersonate a particular resource server.)

- **Clients** We will assume that clients are not trustworthy. Specifically, clients may attempt to obtain tokens that belong to other users and/or modify the tokens issued to them by the authentication server to acquire additional permissions.
- **Other Principals** You should assume that *all* communications in the system are monitored by a *passive adversary*. A passive adversary is able to watch all communication channels in an attempt to learn information, but cannot alter or disrupt any communications in the system. The only form of *active* attack within scope is impersonation of a resource server.

3 Threats to Protect Against

Given the above trust model, we must now consider certain classes of threats that were not addressed in the last phase of the project. In particular, your group must develop defenses against the following classes of threats in this phase of the project:

T1 Unauthorized Token Issuance Due to the fact that clients are untrusted, we must protect against the threat of illegitimate clients requesting tokens from the authentication server. Your implementation must ensure that all clients are authenticated in a secure manner prior to issuing them tokens. That is, we want to ensure that Alice cannot request and receive Bob's security token from the authentication server.

T2 Token Modification/Forgery Users are expected to attempt to modify their tokens to increase their access rights, and to attempt to create forged tokens. Your implementation of tokens must be extended to allow resource servers (or anyone else) to determine whether a token is in fact valid. Specifically, it must be possible for a third-party to verify that a token was in fact issued by a trusted authentication server and was *not* modified after issuance.

T3 Unauthorized Resource Servers The above trust model assumes that properly authenticated resource servers are guaranteed to behave as expected. In order for this guarantee to mean anything, your implementation must ensure that if a user attempts to contact some server, s , then they actually connect to s and not some other server s' .

Note that any user may run a resource server (including modifying the code so that it behaves differently than expected). As such, the authentication server *can not* be required to know about all resource servers. Your mechanism for enabling users to authenticate resource servers should require communication between only the user and the resource server (and possibly client-side application configuration changes).

Hint: You may wish to look into how SSH allows users to authenticate servers.

T4: Information Leakage via Passive Monitoring Since our trust model assumes the existence of passive attackers (e.g., nosy administrators), you must ensure that all communications between your client and server applications are hidden from outside observers. This will ensure that resource contents remain private, and that tokens cannot be stolen in transit.

4 What Do I Need to Do?

This phase of the project has two deliverables. The first deliverable is a semi-formal writeup describing the protection mechanisms that your group proposes to implement, and the second is your actual implementation. We now describe both aspects of the project in greater detail.

4.1 Mechanism Description

The first deliverable for this phase of the project will be a writeup (approximately 3–5 pages) describing the cryptographic mechanisms and protocols that you will implement to address each of the threats identified in Section 3 of this assignment. This writeup should begin with an introductory paragraph or two that broadly surveys the types of cryptographic techniques that your group has decided to use to address threats T1–T4. You should then have one section for each threat, with each section containing the following information:

- Begin by describing the threat treated in this section. This may include describing examples of the threat being exploited by an adversary, a short discussion of why this threat is problematic and needs to be addressed, and/or diagrams showing how the threat might manifest in your group’s current (insecure) implementation.
- Next, provide a short description of the mechanism that you chose to implement to protect against this threat. For interactive protocols, include diagrams explaining the messages exchanged between participating principals. Be sure to explain any cryptographic choices that your group makes: What types of algorithms, modes of operation, and/or key lengths did you choose? **Why?** If shared keys are needed, how are they exchanged? Recall that security is not absolute nor are any tools appropriate for all situations; every component of your design should be intentional and justified relative to the given threat model.
- Lastly, provide a short argument addressing why your proposed mechanism sufficiently addresses this particular threat. This argument should address the correctness of your approach, as well as its overall security. For example, if your mechanism involves a key agreement or key exchange protocol, you should argue that both parties agree on the same key (correctness) and that no other party can figure out the key (security). You do not need a formal proof, but you should convince the reader that an attacker can no longer exploit each threat.

After completing one section for each threat, conclude with a paragraph or two discussing the interplay between your proposed mechanisms, and commenting on the design process that your group followed. Did you discuss other ideas that didn’t pan out before settling on the above-documented approach? Did you end up designing a really interesting protocol suite that addresses multiple threats at once? Use this space to show off your hard work!

Approval of your design: The types of security mechanisms that you will develop during this phase of the project are notoriously finicky and easy to build incorrectly. The goal of this writeup is to focus your group on properly designing and evaluating your mechanisms

before you spend time implementing them. A well-executed writeup will give your group a concrete reference point to discuss, debate, analyze, and (only then!) implement. To encourage such discussion, **10% of your grade for this phase of the project is based upon obtaining the instructor's approval of your design by Nov 10** (the earlier, the better!). You can attach your draft writeup to the **#p3** Discord channel for instructor review. This channel is also a good place to discuss other groups' designs.

Please note that your instructor will not attempt to evaluate a scheme that is described only informally or that lacks the necessary details. If there are flaws in your approaches, you will need to adjust your design and resubmit. This means you may need to submit multiple times, so plan ahead!

4.2 Implementation Requirements

As before, this project will be graded using the CS Linux Cluster. You are responsible for ensuring that your code runs correctly on these machines well before the due date.

You can (and should) use library functions that implement cryptographic primitives, but you should not use any prebuilt protocols or other complex constructions. In particular, avoid external implementations of SSL/TLS/HTTPS, SSH, Kerberos, OTR, IPsec, OAuth, etc. You may use block ciphers, stream ciphers, public key encryption/signing, hash functions, MACs, Diffie-Hellman, secure pseudorandom number generators, and key derivation functions (such as `bcrypt`). If you're unsure about whether you can use a particular library function, ask on **#p3** on Discord.

If your design includes any out-of-band key exchange, be sure to design your system to make this feasible, and justify why it would be realistic in practice (or what should happen instead, if relevant). To ensure a smooth demo, you also need to plan how keys will be exchanged when demonstrating your functionality to your TA. Test run each server and client from a different folder, and make sure that key exchange is being done intentionally and consistently.

5 What (and how) do I submit?

Your grade for this project will be based upon your technical writeup (50%), approval of your design with the instructor on Discord (10%), a demonstration and assessment of the code that your team produces (35%), and scheduling a demo with the TA prior to the deadline (5%).

Within your project repository (your existing `cs1653-project-*` repository from Phase 2), you should include the following files and directories.

- `desc/` In this directory, we provided this project description. No changes are necessary.
- `doc/` In this directory, include all documentation for your project. Include an updated user's manual and technician's guide (separately or combined) for your system based on what you developed in Phase 2. All documentation should be in PDF or HTML format.

In addition, include `doc/phase3-writeup.htm` or `doc/phase3-writeup.pdf`, your writeup that explains your mechanism design as described in Section 4.1.

- `src/` In this directory, include all of your source code that is needed to compile your programs. You may create subdirectories (packages) within this folder if you'd like to better organize the code. Please do not commit any compiled code (e.g., JAR or class files). It is common version-control etiquette not to commit files that can be re-derived from the included source. Also, please do not commit any publicly available libraries that you make use of—include instructions (or, even better, a script) for acquiring those libraries.

Note that we will be evaluating your submission using the CS Linux Cluster as with Phase 2, so you must test your code in that environment prior to the submission deadline!

Each individual's contribution to the group's work will be judged in part by the version control logs. In addition, *each student in your group* will be asked to evaluate the group's performance, including how well everyone is working together and supporting one another.

Your project is due at the precise date and time stated above. We will clone your repository immediately after the due date, so you will be graded on whatever changes have been committed **and pushed** to your repository's main branch by this time. No changes made after this point will be considered in your demo or in grading your project. Make sure your repository is created and you understand the submission process well in advance!